

Shared Memory Programming with OpenMP

Synchronisation



Why is it required?

Recall:

- Need to synchronise actions on shared variables.
- Need to ensure correct ordering of reads and writes.
- Need to protect updates to shared variables (not atomic by default)

BARRIER directive

- No thread can proceed past a barrier until all the other threads have arrived.
- Note that there is an implicit barrier at the end of DO/FOR, SECTIONS and SINGLE directives.

- Syntax:

Fortran: `!$OMP BARRIER`

C/C++: `#pragma omp barrier`

- Either all threads or none must encounter the barrier: otherwise DEADLOCK!!

BARRIER directive (cont)

Example:

```
#pragma omp parallel private(myid,neighb) shared(a,b,c)
{
    myid = omp_get_thread_num();
    neighb = myid - 1;
    if (myid.eq.0) neighb = omp_get_num_threads()-1;
    ...
    a[myid] *= 3.5;
    #pragma omp barrier
    b[myid] = a[neighb] + c;
    ...
}
```

- Barrier required to force synchronisation on **a**

BARRIER directive (cont)

Example:

```
!$OMP PARALLEL PRIVATE(MYID,NEIGHB) SHARED(A,B,C)
  myid = omp_get_thread_num()
  neighb = myid - 1
  if (myid.eq.0) neighb = omp_get_num_threads()-1
  ...
  a(myid) = a(myid)*3.5
!$OMP BARRIER
  b(myid) = a(neighb) + c
  ...
!$OMP END PARALLEL
```

- Barrier required to force synchronisation on **a**

Critical sections



- A critical section is a block of code which can be executed by only one thread at a time.
- Can be used to protect updates to shared variables.

CRITICAL directive

- Syntax:

Fortran: **!\$OMP CRITICAL**

block

!\$OMP END CRITICAL

C/C++: **#pragma omp critical**

structured block

CRITICAL directive (cont)

Example: appending to a shared list

```
#pragma omp parallel for shared(list, N) private(newitem_p)
for (int i=0; i<N; i++) {
    newitem_p = createitem(i);
    #pragma omp critical
    {
        append(&list, p_newitem);
    }
}
```


CRITICAL directive (cont)

Example: appending to a shared list

```
!$OMP PARALLEL DO SHARED(list,n) PRIVATE(newitem)
do i=1,n
    newitem = createitem(i)
    !$OMP CRITICAL
        call append(list,newitem)
    !$OMP END CRITICAL
end do
```

CRITICAL directive (cont)

Example: pushing and popping a task stack

```
#pragma omp parallel shared(stack) private(p_next,p_new,done)
{
while (!done) {
#pragma omp critical
{
    p_next = pop(&stack);
}
    p_new = process(p_next);
#pragma omp critical
{
    if (p_new != NULL) push(p_new,&stack);
    done = isempty(&stack);
}
}
}
```

CRITICAL directive (cont)

Example: pushing and popping a task stack

```
!$OMP PARALLEL SHARED(stack) ,PRIVATE(next,new,done)
    do while (.not. done)
!$OMP CRITICAL
        next = pop(stack)
!$OMP END CRITICAL
        new = process(next)
!$OMP CRITICAL
        if (valid(new)) call push(new,stack)
        done = isempty(stack)
!$OMP END CRITICAL
    end do
!$OMP END PARALLEL
```

ATOMIC directive

- Used to protect a single update to a shared scalar variable of basic type.
- Applies only to a single statement.
- Syntax:

Fortran: **!\$OMP ATOMIC**
statement

where *statement* must have one of these forms:

$x = x \text{ op } \text{expr}$, $x = \text{expr op } x$, $x = \text{intr} (x, \text{expr})$ or
 $x = \text{intr} (\text{expr}, x)$

op is one of **+**, *****, **-**, **/**, **.and.**, **.or.**, **.eqv.**, or **.neqv.**

intr is one of **MAX**, **MIN**, **IAND**, **IOR** or **IEOR**

ATOMIC directive (cont)

C/C++: `#pragma omp atomic`
statement

where *statement* must have one of the forms:

x binop = expr, *x++*, *++x*, *x--*, or *--x*


and *binop* is one of *+*, ***, *-*, */*, *&*, *^*, *<<*, or *>>*

- Note that the evaluation of *expr* is not atomic.
- May be more efficient than using CRITICAL directives, e.g. if different array elements can be protected separately.
- No interaction with CRITICAL directives

ATOMIC directive (cont)

Example (compute degree of each vertex in a graph):

```
#pragma omp parallel for
    for (j=0; j<nedges; j++){
#pragma omp atomic
        degree[edge[j].vertex1]++;
#pragma omp atomic
        degree[edge[j].vertex2]++;
    }
```



Lock routines

- Occasionally we may require more flexibility than is provided by CRITICAL directive.
- A lock is a special variable that may be *set* by a thread. No other thread may *set* the lock until the thread which set the lock has *unset* it.
- Setting a lock can either be blocking or non-blocking.
- A lock must be initialised before it is used, and may be destroyed when it is no longer required.
- Lock variables should not be used for any other purpose.

Lock routines - syntax

Fortran:

```
USE OMP_LIB  
  
SUBROUTINE OMP_INIT_LOCK(OMP_LOCK_KIND var)  
SUBROUTINE OMP_SET_LOCK(OMP_LOCK_KIND var)  
LOGICAL FUNCTION OMP_TEST_LOCK(OMP_LOCK_KIND var)  
SUBROUTINE OMP_UNSET_LOCK(OMP_LOCK_KIND var)  
SUBROUTINE OMP_DESTROY_LOCK(OMP_LOCK_KIND var)
```

var should be an INTEGER of the same size as addresses (e.g. INTEGER*8 on a 64-bit machine)

OMP_LIB defines OMP_LOCK_KIND

Lock routines - syntax

C/C++:

```
#include <omp.h>

void omp_init_lock(omp_lock_t *lock);
void omp_set_lock(omp_lock_t *lock);
int omp_test_lock(omp_lock_t *lock);
void omp_unset_lock(omp_lock_t *lock);
void omp_destroy_lock(omp_lock_t *lock);
```


Lock example

Example (compute degree of each vertex in a graph):

```
omp_lock_t lockvar[nvertices];

for (i=0; i<nvertexes; i++){
    omp_init_lock(&lockvar[i]);
}

#pragma omp parallel for
    for (j=0; j<nedges; j++){
        omp_set_lock(&lockvar[edge[j].vertex1]);
        degree[edge[j].vertex1]++;
        omp_unset_lock(&lockvar[edge[j].vertex1]);
        omp_set_lock(&lockvar[edge[j].vertex2]);
        degree[edge[j].vertex2]++;
        omp_unset_lock(&lockvar[edge[j].vertex2]);
    }
```



Exercise: Molecular dynamics



- The code supplied is a simple molecular dynamics simulation of the melting of solid argon.
- Computation is dominated by the calculation of force pairs in subroutine **forces**.
- Parallelise this routine using a DO/FOR directive and critical sections.
 - Watch out for PRIVATE and REDUCTION variables.
 - Choose a suitable loop schedule
- Extra exercise: can you improve the performance by using locks, or atomics, or by using a reduction array.

Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

<https://creativecommons.org/licenses/by-nc-sa/4.0/>

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Acknowledge EPCC as follows: “© EPCC, The University of Edinburgh, www.epcc.ed.ac.uk”

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.